

Exercise sheet

Lecturer: Georgios P. Karagiannis

georgios.karagiannis@durham.ac.uk

Part 1. Elements of convex learning problems

Exercise 1. (★) Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ such that $f(w) = g(\langle w, x \rangle + y)$ for some $x \in \mathbb{R}^d, y \in \mathbb{R}$. Show that: If g is convex function then f is convex function.

Solution. Let $u, v \in \mathbb{R}^d$ and $\alpha \in [0, 1]$. It is

$$\begin{aligned}
 f(\alpha u + (1 - \alpha)v) &= g(\langle \alpha u + (1 - \alpha)v, x \rangle + y) \\
 &= g(\langle \alpha u, x \rangle + \langle (1 - \alpha)v, x \rangle + y) \\
 &= g(\alpha \langle u, x \rangle + y + (1 - \alpha)(\langle v, x \rangle + y)) & y = \alpha y + (1 - \alpha)y \\
 &\leq \alpha g(\langle u, x \rangle + y) + (1 - \alpha)g(\langle v, x \rangle + y) & (g \text{ is convex}) \\
 &= \alpha f(u) + (1 - \alpha)f(v)
 \end{aligned}$$

Exercise 2. (★) Let functions g_1 be ρ_1 -Lipschitz and g_2 be ρ_2 -Lipschitz. Then, show that, f with $f(x) = g_1(g_2(x))$ is $\rho_1\rho_2$ -Lipschitz.

Solution.

$$\begin{aligned}
 |f(w_1) - f(w_2)| &= |g_1(g_2(w_1)) - g_1(g_2(w_2))| \\
 &\leq \rho_1 |g_2(w_1) - g_2(w_2)| \\
 &\leq \rho_1 \rho_2 |w_1 - w_2|
 \end{aligned}$$

Exercise 3. (★) Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ with $f(w) = g(\langle w, x \rangle + y)$ $x \in \mathbb{R}^d$ and $y \in \mathbb{R}$. Let $g : \mathbb{R} \rightarrow \mathbb{R}$ be a β -smooth function. Then show that f is a $(\beta \|x\|^2)$ -smooth.

Hint:: You may use Cauchy-Schwarz inequality $\langle y, x \rangle \leq \|y\| \|x\|$

$$f(v) = g(\langle w, x \rangle + y)$$

$$\leq g(\langle w, x \rangle + y) + g'(\langle w, x \rangle + y) \langle v - w, x \rangle + \frac{\beta}{2} (\langle v - w, x \rangle)^2 \quad (g \text{ is smooth})$$

$$\leq g(\langle w, x \rangle + y) + g'(\langle w, x \rangle + y) \langle v - w, x \rangle + \frac{\beta}{2} (\|v - w\| \|x\|)^2 \quad (\text{Cauchy-Schwarz inequality})$$

$$= f(w) + \langle \nabla f(w), v - w \rangle + \frac{\beta \|x\|^2}{2} \|v - w\|^2$$

Exercise 4. (★) Show that $f : S \rightarrow \mathbb{R}$ is ρ -Lipschitz over an open convex set S if and only if for all $w \in S$ and $v \in \partial f(w)$ it is $\|v\| \leq \rho$.

Hint:: You may use Cauchy-Schwarz inequality $\langle y, x \rangle \leq \|y\| \|x\|$

Solution. \Rightarrow Let $f : S \rightarrow \mathbb{R}$ be ρ -Lipschitz over convex set S , $w \in S$ and $v \in \partial f(w)$.

- Since S is open we get that there exist $\epsilon > 0$ such as $u := w + \epsilon \frac{v}{\|v\|}$ where $u \in S$. So $\langle u - w, v \rangle = \epsilon \|v\|$ and $\|u - w\| = \epsilon$.
- From the subgradient definition we get

$$f(u) - f(w) \geq \langle u - w, v \rangle = \epsilon \|v\|$$

- From the Lipschitzness of $f(\cdot)$ we get

$$f(u) - f(w) \leq \rho \|u - w\| = \rho \epsilon$$

Therefore $\|v\| \leq \rho$.

\Leftarrow It is for all $w \in S$ and $v \in \partial f(w)$ it is $\|v\| \leq \rho$.

- For any $u \in S$, it is

$$\begin{aligned} f(w) - f(u) &\leq \langle v, w - u \rangle && (\text{because } v \in \partial f(w)) \\ (1) \quad &\leq \|v\| \|w - u\| && \text{by Cauchy-Schwarz inequality} \\ &\leq \rho \|w - u\| && \text{because } \|v\| \leq \rho \end{aligned}$$

- Similarly it results $u, w \in S$

$$f(w) - f(u) \leq \langle v, u - w \rangle \|v\| \leq \|v\| \|u - w\| \leq \rho \|u - w\|$$

from (1) because w, u can be swapped in (1) as they both are any values in S .

Exercise 5. (★) Let $g_1(w), \dots, g_r(w)$ be r convex functions, and let $f(\cdot) = \max_{\forall j} (g_j(\cdot))$. Show that for some w it is $\nabla g_k(w) \in \partial f(w)$ where $k = \arg \max_j (g_j(w))$ is the index of function $g_j(\cdot)$ presenting the greatest value at w .

Solution. Since g_k is convex, for all u

$$g_k(u) \geq g_k(w) + \langle u - w, \nabla g_k(w) \rangle$$

However $f(u) = \max_{\forall j} (g_j(u)) \geq g_k(u)$ for any j , and $f(w) = g_k(w)$ at w . Then

$$\begin{aligned} f(u) &\geq g_k(u) \\ &\geq g_k(w) + \langle u - w, \nabla g_k(w) \rangle \\ &= f(w) + \langle u - w, \nabla g_k(w) \rangle \end{aligned}$$

Then by the definition of the sub-gradient $\nabla g_k(w) \in \partial f(w)$

Exercise 6. (★) Consider the regression learning problem $(\mathcal{H}, \mathcal{Z}, \ell)$ with predictor rule $h(x) = \langle w, x \rangle$ labeled by some unknown parameter $w \in \mathcal{W}$, loss function $\ell(w, (x, y)) = (\langle w, x \rangle - y)^2$, feature $x \in \mathcal{X}$, and target $y \in \mathbb{R}$. Let $\mathcal{W} = \mathcal{X} = \{\omega \in \mathbb{R}^d : |\omega| \leq \rho\}$ for some $\rho > 0$.

- (1) Show that the resulting learning problem is Convex-Lipschitz-Bounded learning problem.
- (2) Specify the parameters of Lipschitzness.

Solution. According to the definitions given in the lecture:

- Convex-Lipschitz-Bounded Learning Problem $(\mathcal{H}, \mathcal{Z}, \ell)$ with parameters ρ , and B , is called the learning problem whose the hypothesis class \mathcal{H} is a convex set, for all $w \in \mathcal{H}$ it is $\|w\| \leq B$, and the loss function $\ell(\cdot, z)$ is convex and ρ -Lipschitz function for all $z \in \mathcal{Z}$.

I have:

Convexity: The function $g : \mathbb{R} \rightarrow \mathbb{R}$, defined by $g(a) = a^2$ is convex. Eg. $\frac{d^2}{da^2} g(a) = 2 \geq 0$ is non-negative. The convexity of $\ell(w, z = (x, y))$ for all z follows as a composition of g with a linear function.

Lipschitzness: The function $g(a) = a^2$ is 1-Lipschitz since It is

$$|g(a_2) - g(a_1)| = |a_2^2 - a_1^2| = |(a_2 + a_1)(a_2 - a_1)| \leq 2\rho(a_2 - a_1) = 2\rho|a_2 - a_1|$$

Hence because $|x| \leq \rho$, $g(a)$ is $2\rho^2$ -Lipschitz as a composition.

Boundness: The norm of each hypothesis w is bounded by ρ according to the assumptions.

Therefore,

- (1) the learning problem under consideration is a Convex-Lipschitz-Bounded learning problem.
- (2) the parameter of Lipschitzness is $2\rho^2$.

Exercise 7. (★) If f is λ -strongly convex and u is a minimizer of f then for any w

$$f(w) - f(u) \geq \frac{\lambda}{2} \|w - u\|^2$$

Hint:: Use the definition, and set $\alpha \rightarrow 0$.

Solution.

The following is given as a homework (Formative assessment 1)

Exercise 8. (★) Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a convex and β -smooth function.

(1) Show that for $v, w \in \mathbb{R}^d$

$$f(v) - f(w) \in \left(\langle \nabla f(w), v - w \rangle, \langle \nabla f(w), v - w \rangle + \frac{\beta}{2} \|v - w\|^2 \right)$$

(2) Show that for $v, w \in \mathbb{R}^d$ such that $v = w - \frac{1}{\beta} \nabla f(w)$, it is

$$\frac{1}{2\beta} \|\nabla f(w)\|^2 \leq f(w) - f(v)$$

(3) Additionally assume that $f(x) > 0$ for all $x \in \mathbb{R}^d$. Show that for $w \in \mathbb{R}^d$,

$$\|\nabla f(w)\| \leq \sqrt{2\beta f(w)}$$

Solution.

(1) If $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is β -smooth then it is

$$f(v) \leq f(w) + \langle \nabla f(w), v - w \rangle + \frac{\beta}{2} \|v - w\|^2$$

$$f(v) - f(w) \leq \langle \nabla f(w), v - w \rangle + \frac{\beta}{2} \|v - w\|^2$$

If it is convex then it is

$$f(v) \geq f(w) + \langle \nabla f(w), v - w \rangle$$

$$f(v) - f(w) \geq \langle \nabla f(w), v - w \rangle$$

Together these conditions imply upper and lower bounds

$$f(v) - f(w) \in \left(\langle \nabla f(w), v - w \rangle, \langle \nabla f(w), v - w \rangle + \frac{\beta}{2} \|v - w\|^2 \right)$$

(2) For $v, w \in \mathbb{R}^d$ such that $v = w - \frac{1}{\beta} \nabla f(w)$, it is

$$f(v) \leq f(w) + \langle \nabla f(w), v - w \rangle + \frac{\beta}{2} \|v - w\|_2^2 \quad (\text{due to smoothness})$$

$$\iff f(w) - f(v) \leq f(w) - f(v)$$

$$\iff \langle \nabla f(w), v - w \rangle + \frac{\beta}{2} \|v - w\|_2^2 \leq f(w) - f(v)$$

$$\iff \left\langle \nabla f(w), \frac{1}{\beta} \nabla f(w) \right\rangle + \frac{\beta}{2} \left\| \frac{1}{\beta} \nabla f(w) \right\|_2^2 \leq f(w) - f(v)$$

$$\iff \frac{1}{2\beta} \|\nabla f(w)\|^2 \leq f(w) - f(v)$$

$$\|\nabla f(w)\|^2 \leq 2\beta (f(w) - f(v))$$

as $f(\cdot) \geq 0$

$$\|\nabla f(w)\|^2 \leq 2\beta f(w)$$

(3) From part 2, this is obvious because $f(x) > 0$ for all $x \in \mathbb{R}^d$, as

$$\|\nabla f(w)\|^2 \leq 2\beta f(w) \Leftrightarrow \|\nabla f(w)\| \leq \sqrt{2\beta f(w)}$$

The following is given as a homework (Formative assessment 1)

Exercise 9. (★) Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a λ -strongly convex function. Assume that w^* is a minimizer of f i.e.

$$w^* = \arg \min_w \{f(w)\}$$

Show that for any $w \in \mathbb{R}^d$ it holds

$$f(w) - f(w^*) \geq \frac{\lambda}{2} \|w - w^*\|^2$$

Hint: Use the definition of λ -strongly convex function, properly rearrange it, and ...

Solution. We use the definition of λ -strongly convex function; i.e. for all w, u , and $\alpha \in (0, 1)$ we have

$$\begin{aligned} f(aw + (1 - \alpha)u) &\leq af(w) + (1 - \alpha)f(u) - \frac{\lambda}{2}\alpha(1 - \alpha)\|w - u\|^2 \Leftrightarrow \\ \frac{f(aw + (1 - \alpha)u) - f(u)}{\alpha} &\leq f(w) + f(u) - \frac{\lambda}{2}(1 - \alpha)\|w - u\|^2 \end{aligned}$$

For $u = w^*$ it is

$$\frac{f(aw + (1 - \alpha)w^*) - f(w^*)}{\alpha} \leq f(w) + f(w^*) - \frac{\lambda}{2}(1 - \alpha)\|w - w^*\|^2$$

When $a \rightarrow 0$

$$\frac{\lambda}{2}\alpha(1 - \alpha)\|w - w^*\|^2 \rightarrow 0$$

I know that w^* is the minimizer of f . So 0 is the minimizer of g with $g(a) = f(aw + (1 - \alpha)w^*)$ hence when $a \rightarrow 0$

$$\frac{f(aw + (1 - \alpha)w^*) - f(w^*)}{\alpha} \rightarrow \left. \frac{d}{da} g(a) \right|_{a=0}$$

So

$$0 \leq f(w) + f(w^*) - \frac{\lambda}{2}\|w - w^*\|^2$$

which concludes the proof.

Exercise 10. (★) Show that the function $J(x; \lambda) = \lambda \|x\|^2$ is 2λ -strongly convex

Solution. We just need to check that for all w, u , and $\alpha \in (0, 1)$ we have

$$\begin{aligned} J(aw + (1 - \alpha)u; \lambda) &\leq aJ(w; \lambda) + (1 - \alpha)J(u; \lambda) - \frac{2\lambda}{2}\alpha(1 - \alpha)\|w - u\|^2 \Leftrightarrow \\ \|aw + (1 - \alpha)u\|_2^2 &\leq a\|w\|_2^2 + (1 - \alpha)\|u\|_2^2 - a(1 - \alpha)\|w - u\|_2^2 \Leftrightarrow 0 \leq 0 \end{aligned}$$

Part 2. Stochastic learning

Exercise 11. (★) Assume a Bayesian model

$$\begin{cases} z_i|w & \stackrel{\text{ind}}{\sim} f(z_i|w), \quad i = 1, \dots, n \\ w & \sim f(w) \end{cases}$$

and consider that our objective is the discovery of MAP estimate w^* i.e.

$$w^* = \arg \min_{w \in \Theta} (-\log(L_n(w)) - f(w)) = \arg \min_{w \in \Theta} \left(-\sum_{i=1}^n \log(f(z_i|w)) - \log(f(w)) \right)$$

by using SGD with update

$$w^{(t+1)} = w^{(t)} + \eta_t \left(\frac{n}{m} \sum_{j \in \mathcal{J}^{(t)}} \nabla_w \log(f(z_j|w^{(t)})) + \nabla_w \log(f(w^{(t)})) \right)$$

for some randomly selected set $\mathcal{J}^{(t)} \subseteq \{1, \dots, n\}^m$ of m integers from 1 to n via simple random sampling (SRS) with replacement. Show that

$$\mathbb{E}_{\mathcal{J}^{(t)} \sim \text{simple-random-sampling}} \left(\frac{n}{m} \sum_{j \in \mathcal{J}^{(t)}} \nabla_w \log(f(z_j|w^{(t)})) \right) = \sum_{i=1}^n \nabla_w \log(f(z_i|w^{(t)}))$$

Solution. It is

$$\begin{aligned} \mathbb{E}_{\mathcal{J}^{(t)} \sim \text{SRS}} \left(\frac{n}{m} \sum_{j \in \mathcal{J}^{(t)}} \nabla_w \log(f(z_j|w^{(t)})) \right) &= \frac{n}{m} \sum_{j \in \mathcal{J}^{(t)}} \mathbb{E}_{\mathcal{J}^{(t)} \sim \text{SRS}} \left(\nabla_w \log(f(z_j|w^{(t)})) \right) \\ &= \frac{n}{m} \sum_{j \in \mathcal{J}^{(t)}} \mathbb{E}_{\mathcal{J}^{(t)} \sim \text{SRS}} \left(\nabla_w \log(f(z_j|w^{(t)})) \right) \\ &= \frac{n}{m} \sum_{j \in \mathcal{J}^{(t)}} \frac{1}{n} \sum_{i=1}^n \nabla_w \log(f(z_i|w^{(t)})) \\ &= \sum_{i=1}^n \nabla_w \log(f(z_i|w^{(t)})) \end{aligned}$$

It is $\mathbb{E}_{\mathcal{J}^{(t)} \sim \text{SRS}} (\nabla_w \log(f(z_j|w^{(t)}))) = \frac{1}{n} \sum_{i=1}^n \nabla_w \log(f(z_i|w^{(t)}))$ because the expectation is under the probability I get randomly an integer and for the j th on the probability is $1/n$ due to the random scheme. Also $|\mathcal{J}^{(t)}| = m$.

Exercise 12. (★) Let $\{v_t; t = 1, \dots, T\}$ be a sequence of vectors with $v_t \in \mathbb{R}^d$ and $d \in \mathbb{N} - \{0\}$. Consider an algorithm producing $\{w^{(t)}; t = 1, 2, 3, \dots\}$ with

$$\begin{aligned} w^{(1)} &= 0 \\ w^{(t+1)} &= w^{(t)} - \eta v_t \end{aligned}$$

$w_t \in \mathbb{R}^d$ and $d \in \mathbb{N} - \{0\}$. Show that

(1) it is

$$\langle w^{(t)} - w^*, v_t \rangle = \frac{1}{2\eta} \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \|v_t\|^2$$

Hint:: Recall that

$$\|x + y\|_2^2 = \|x\|_2^2 + \|y\|_2^2 + 2\langle x, y \rangle, \quad \forall x, y \in \mathbb{R}^d, d \in \mathbb{N} - \{0\}$$

(2) it is

$$\sum_{t=1}^T \langle w^{(t)} - w^*, v_t \rangle = \frac{1}{2\eta} \sum_{t=1}^T \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2$$

(3) (continue) it is

$$\sum_{t=1}^T \langle w^{(t)} - w^*, v_t \rangle \leq \frac{\|w^*\|^2}{2\eta} + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2$$

Solution.

(1) It is

$$\begin{aligned} \langle w^{(t)} - w^*, v_t \rangle &= \frac{1}{\eta} \langle w^{(t)} - w^*, \eta v_t \rangle \\ &= \frac{1}{\eta} \left(-\langle w^{(t)} - w^*, -\eta v_t \rangle \right) \end{aligned}$$

Then by using the Hint as

$$\langle x, y \rangle = \frac{1}{2} \left(\|x + y\|_2^2 - \|x\|_2^2 - \|y\|_2^2 \right)$$

for $x = w^{(t)} - w^* \in \mathbb{R}^d$ and $y = -\eta v_t \in \mathbb{R}^d$, I get

$$\begin{aligned} \langle w^{(t)} - w^*, v_t \rangle &= \frac{1}{2\eta} \left(-\|w^{(t)} - w^* - \eta v_t\|^2 + \|w^{(t)} - w^*\|^2 + \|\eta v_t\|^2 \right) \\ &= \frac{1}{2\eta} \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 + \eta^2 \|v_t\|^2 \right) \\ &= \frac{1}{2\eta} \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \|v_t\|^2 \end{aligned}$$

(2) So

$$\begin{aligned} \sum_{t=1}^T \langle w^{(t)} - w^*, v_t \rangle &= \frac{1}{2\eta} \sum_{t=1}^T \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2 \\ &= \frac{1}{2\eta} \left(\|w^{(1)} - w^*\|^2 - \|w^{(T+1)} - w^*\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2 \end{aligned}$$

(3) So

$$\begin{aligned}
\sum_{t=1}^T \langle w^{(t)} - w^*, v_t \rangle &= \frac{1}{2\eta} \left(\|w^{(1)} - w^*\|^2 - \|w^{(T+1)} - w^*\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2 \\
&\leq \frac{1}{2\eta} \|w^{(1)} - w^*\|^2 + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2 \\
&= \frac{1}{2\eta} \|w^*\|^2 + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2
\end{aligned}$$

Exercise 13. (★) Let $\{v_t; t = 1, \dots, T\}$ be a sequence of vectors. Consider an algorithm producing $\{w^{(t)}; t = 1, 2, 3, \dots\}$ with

$$\begin{aligned}
w^{(1)} &= 0 \\
w^{(t+\frac{1}{2})} &= w^{(t)} - \eta v_t \\
w^{(t+1)} &= \arg \min_{w \in \mathcal{H}} \left(\|w - w^{(t+\frac{1}{2})}\| \right)
\end{aligned}$$

for $t = 1, \dots, T$.

Hint: You can use the following Lemma

(Projection Lemma): Let \mathcal{H} be a closed convex set and let v be the projection of w onto \mathcal{H} , i.e.

$$v = \arg \min_{x \in \mathcal{H}} \|x - w\|^2$$

then for every $u \in \mathcal{H}$ it is

$$\|v - u\|^2 \leq \|w - u\|^2$$

Show that

(1) it is

$$\langle w^{(t)} - w^*, v_t \rangle \leq \frac{1}{2\eta} \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \|v_t\|^2$$

(2) it is

$$\sum_{t=1}^T \langle w^{(t)} - w^*, v_t \rangle \leq \frac{1}{2\eta} \sum_{t=1}^T \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2$$

(3) (continue) it is

$$\sum_{t=1}^T \langle w^{(t)} - w^*, v_t \rangle \leq \frac{\|w^*\|^2}{2\eta} + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2$$

Comment: Above we show that Lemma ?? from “Handout ??: Gradient descent” holds even when a projection step is included. Hence, even if a projection step is included after the update step of the recursion of GD algorithm or the SGD algorithm the analysis in Section ?? in “Handout ??: Gradient descent” holds. Hence, even if a projection step is included

after the update step of the recursion of SGD algorithm or the SGD algorithm the analysis in Section ?? in “Handout ??: Stochastic gradient descent” holds.

Solution.

(1) It is

$$\begin{aligned}
\langle w^{(t)} - w^*, v_t \rangle &= \frac{1}{\eta} \langle w^{(t)} - w^*, \eta v_t \rangle \\
&= \frac{1}{2\eta} \left(-\|w^{(t)} - w^* - \eta v_t\|^2 + \|w^{(t)} - w^*\|^2 + \eta^2 \|v_t\|^2 \right) \\
&= \frac{1}{2\eta} \left(-\|w^{(t+\frac{1}{2})} - w^*\|^2 + \|w^{(t)} - w^*\|^2 + \eta^2 \|v_t\|^2 \right) \\
&= \frac{1}{2\eta} \left(-\|w^{(t+\frac{1}{2})} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \|v_t\|^2 \\
&\leq \frac{1}{2\eta} \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \|v_t\|^2
\end{aligned}$$

because from the Projection Lemma

$$\|w^{(t+1)} - w^*\|^2 \leq \|w^{(t+\frac{1}{2})} - w^*\|^2$$

(2) So

$$\begin{aligned}
\sum_{t=1}^T \langle w^{(t)} - w^*, v_t \rangle &\leq \frac{1}{2\eta} \sum_{t=1}^T \left(-\|w^{(t+1)} - w^*\|^2 + \|w^{(t)} - w^*\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2 \\
&= \frac{1}{2\eta} \left(\|w^{(1)} - w^*\|^2 - \|w^{(T+1)} - w^*\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2
\end{aligned}$$

(3) So

$$\begin{aligned}
\sum_{t=1}^T \langle w^{(t)} - w^*, v_t \rangle &\leq \frac{1}{2\eta} \left(\|w^{(1)} - w^*\|^2 - \|w^{(T+1)} - w^*\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2 \\
&\leq \frac{1}{2\eta} \|w^{(1)} - w^*\|^2 + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2 \\
&= \frac{1}{2\eta} \|w^*\|^2 + \frac{\eta}{2} \sum_{t=1}^T \|v_t\|^2
\end{aligned}$$

The following is given as a homework (Formative assessment 2)

Exercise 14. (★) ¹Consider the binary classification problem with inputs $x \in \mathcal{X}$ where $\mathcal{X} := \{x \in \mathbb{R}^d : \|x\|_2 \leq L\}$ for some given value $L > 0$, target $y \in \mathcal{Y}$ where $\mathcal{Y} := \{-1, +1\}$, and prediction

¹We use standard notation

$$\text{sign}(\xi) = \begin{cases} -1, & \text{if } \xi < 0 \\ +1, & \text{if } \xi > 0 \end{cases}$$

± 1 means either -1 or $+1$, $\mathbb{R}_+ := (0, +\infty)$, and $\|x\|_2 := \sqrt{\sum_{j=1}^d (x_j)^2}$ for the Euclidean distance.

rule $h_w : \mathbb{R}^d \rightarrow \{-1, +1\}$ with

$$(2) \quad h_w(x) = \text{sign}(w^\top x)$$

$$(3) \quad = \text{sign}\left(\sum_{j=1}^d w_j x_j\right)$$

Let the hypothesis class is

$$(4) \quad \mathcal{H} = \left\{x \rightarrow w^\top x : \forall w \in \mathbb{R}^d\right\}$$

In other words, the hypothesis $h_w \in \mathcal{H}$ is parametrized by $w \in \mathbb{R}^d$, it receives an input vector $x \in \mathcal{X} := \mathbb{R}^d$ and it returns the label $y = \text{sign}(w^\top x) \in \mathcal{Y} := \{\pm 1\}$ where

$$\text{sign}(\xi) = \begin{cases} -1, & \text{if } \xi < 0 \\ +1, & \text{if } \xi > 0 \end{cases}$$

Consider a loss function $\ell : \mathbb{R}^d \rightarrow \mathbb{R}_+$ with

$$(5) \quad \ell(w, z = (x, y)) = \max(0, 1 - yw^\top x) + \lambda \|w\|_2^2$$

for some given value $\lambda > 0$.

Assume there is available a dataset of examples $S_n = \{z_i = (x_i, y_i) ; i = 1, \dots, n\}$ of size n .

Do the following:

- (1) Show that the function $f : \mathbb{R} \rightarrow \mathbb{R}_+$ with $f(x) = \max(0, 1 - x)$ is convex in \mathbb{R} ; and show that the loss (5) is convex.

Hint:: You may use Proposition ?? from Handout ?? : Elements of convex learning problems.

- (2) Show that the loss $\ell(w, z)$ for $\lambda = 0$ (5) is L -Lipschitz (with respect to w) when $x \in \mathcal{X}$ where $\mathcal{X} := \{x \in \mathbb{R}^d : \|x\|_2 \leq L\}$.

Hint:: You may use the definition of Lipschitz function. Without loss of generality, you can consider any $w_1 \in \mathbb{R}^d$ and $w_2 \in \mathbb{R}^d$ such that $1 - yw_2^\top x \leq 1 - yw_1^\top x$, and then take cases $1 - yw_2^\top x > \text{or} < 0$ and $1 - yw_1^\top x > \text{or} < 0$ to deal with the max.

- (3) Construct the set of sub-gradients $\partial f(x)$ for $x \in \mathbb{R}$ of the function $f : \mathbb{R} \rightarrow \mathbb{R}_+$ with $f(x) = \max(0, 1 - x)$. Show that the vector v with

$$v = \begin{cases} 2\lambda w, & yw^\top x > 1 \\ 2\lambda w, & yw^\top x = 1 \\ -yx + 2\lambda w, & yw^\top x < 1 \end{cases}$$

is $v \in \partial_w \ell(w, z = (x, y))$, aka a sub-gradient of $\ell(w, z = (x, y))$ at w , for any $w \in \mathbb{R}^d$.

- (4) Write down the algorithm of online AdaGrad (Adaptive Stochastic Gradient Descent) with learning rate $\eta_t > 0$, batch size m , and termination criterion $t > T_{\max}$ for some $T_{\max} > 0$ in

order to discover w^* such as

$$(6) \quad w^* = \arg \min_{\forall w: h_w \in \mathcal{H}} (\mathbb{E}_{z \sim g} (\ell(w, z = (x, y))))$$

The formulas in your algorithm should be implemented for the above learning problem and tailored to 2, 4, and 5.

- (5) Use the R code given below in order to generate the dataset of observed examples $S_n = \{z_i = (x_i, y_i)\}_{i=1}^n$ that contains $n = 10^6$ examples with inputs x of dimension $d = 2$. Consider $\lambda = 0$. Use a seed $w^{(0)} = (0, 0)^\top$.
 - (a) By using appropriate values for m , η_t and T_{\max} , code in R the algorithm you designed in part 4, and run it.
 - (b) Plot the trace plots for each of the dimensions of the generated chain $\{w^{(t)}\}$ against the iteration t .
 - (c) Report the value of the output w_{adaGrad}^* (any type) of the algorithm as the solution to (6).
 - (d) To which cluster y (i.e., -1 or 1) $x_{\text{new}} = (1, 0)^\top$ belongs?

```

# R code. Run it before you run anything else
#
data_generating_model <- function(n,w) {
  z <- rep( NaN, times=n*3 )
  z <- matrix(z, nrow = n, ncol = 3)
  z[,1] <- rep(1,times=n)
  z[,2] <- runif(n, min = -10, max = 10)
  p <- w[1]*z[,1] + w[2]*z[,2] p <- exp(p) / (1+exp(p))
  z[,3] <- rbinom(n, size = 1, prob = p)
  ind <- (z[,3]==0)
  z[ind,3] <- -1
  x <- z[,1:2]
  y <- z[,3]
  return(list(z=z, x=x, y=y))
}
n_obs <- 1000000
w_true <- c(-3,4)
set.seed(2023)
out <- data_generating_model(n = n_obs, w = w_true)
set.seed(0)
z_obs <- out$z #z=(x,y)
x <- out$x
y <- out$y
#z_obs2=z_obs
#z_obs2[z_obs[,3]==-1,3]=0
#w_true <- as.numeric(glm(z_obs2[,3]~ 1+ z_obs2[,2],family = "binomial"
)$coefficients)

```

Solution.

- (1) $f_1(x) = 0$ is convex, $f_2(x) = 1 - x$ is convex, hence from the example in Handout 1, $f(x) = \max(f_1(x), f_2(x))$ is convex as well. Regarding the loss function, we just have $f_2(w) = 1 - yx^\top w$ which is convex as a composition due to linearity.
- (2) Given a fixed example $(x, y) \in \{x \in \mathbb{R}^d : \|x'\|_2 \leq R\} \times \{-1, 1\}$.

Assume $w_1, w_2 \in \mathbb{R}^d$. Let $\ell_i = \max\{0, 1 - yx^\top w_i\}$, for $i = 1, 2$. It suffices to show that $|\ell_1 - \ell_2|_2 \leq R|w_1 - w_2|_2$. I take cases

Case-1: Assume $yx^\top w_1 \geq 1$ and $yx^\top w_2 \geq 1$ then $|\ell_1 - \ell_2|_2 = 0 \leq R|w_1 - w_2|_2$

Case-2: Assume that at least one of $yx^\top w_1 < 1$ or $yx^\top w_2 < 1$ but not both is true.

Assume without loss of generality that $1 - yx^\top w_1 < 1 - yx^\top w_2$. Then

$$\begin{aligned}
|\ell_1 - \ell_2|_2 &= \ell_1 - \ell_2 \\
&= 1 - yx^\top w_1 - \max(0, 1 - yx^\top w_2) \\
&\leq 1 - yx^\top w_1 - (1 - yx^\top w_2) \\
&= yx^\top (w_2 - w_1) \\
&\leq y \left\| x^\top \right\|_2 \|w_1 - w_2\|_2 \quad \text{because } a^\top b \leq \|a\| \|b\|
\end{aligned}$$

(3) It is

$$f(x) = \max(0, 1 - x) = \begin{cases} 0 & x > 1 \\ 0 & x = 1 \\ 1 - x & x < 1 \end{cases}$$

- For $x > 1$, f is differentiable so $\partial f(x) = \{f'(x)\} = \{0\}$.
- For $x < 1$, f is differentiable so $\partial f(x) = \{f'(x)\} = \{-1\}$.
- For $x = 1$, f is not differentiable. By definition I have that v is subgradient of $f(x)$ at $x = 0 \in S$ if

$$\forall u \in \mathbb{R}, \quad f(u) \geq f(x) + \langle u - x, v \rangle$$

So, for $u \geq 1$, it is $0 \geq (u - 1)v \implies v \leq 0$, and for $u < 1$ it is $(1 - u) \geq (u - 1)v \implies v \geq -1$. Hence the common space is $v \in [0, 1]$ So $\partial f(x) = [0, 1]$. Hence,

$$\partial f(x) = \begin{cases} 0, & x > 1 \\ [-1, 0], & x = 1 \\ -1, & x < 1 \end{cases}$$

Now regarding the loss $\partial_w \ell(w, z = (x, y))$

- for $yw^\top x > 1$ it is differentiable so $\nabla_w \ell(w, z = (x, y)) = \nabla_w (0 + \lambda \sum_{j=1}^d w_j^2) = 2\lambda w$;
as

$$\frac{d}{dw_j} \sum_{j'=1}^d w_{j'}^2 = 2\lambda w_j$$

- for $yw^\top x < 1$ it is differentiable so $\nabla_w \ell(w, z = (x, y)) = \nabla_w (1 - yw^\top x + \lambda \sum_{j=1}^d w_j^2) = yx + 2\lambda w$ as

$$\frac{d}{dw_j} (1 - yw^\top x) = \frac{d}{dw_j} \left(1 - y \sum_{j'=1}^d w_{j'} x_{j'} \right) = -yx_j$$

- for $yw^\top x = 1$, $v = 0$ satisfies the definition of the sub-gradient

$$\begin{aligned} \forall u, \quad f(u) &\geq \cancel{f(w)}^0 + \langle u - w, v \rangle \\ \max(0, 1 - yu^\top x) &\geq 0 + (u - w)^\top 0 \end{aligned}$$

So

$$\begin{aligned} \partial \ell(w, z = (x, y)) &= \partial \left(\max(0, 1 - yw^\top x) + \lambda \|w\|_2^2 \right) \\ &= \partial \left(\max(0, 1 - yw^\top x) \right) + \partial \left(\lambda \|w\|_2^2 \right) \\ &= \partial \left(\max(0, 1 - yw^\top x) \right) + \nabla \left(\lambda \|w\|_2^2 \right) \\ &= 0 + 2\lambda w \end{aligned}$$

but $\partial \left(\lambda \|w\|_2^2 \right) = \left\{ \nabla \left(\lambda \|w\|_2^2 \right) \right\}$ because $\lambda \|w\|_2^2$ is differentiable. Hence

$$\partial \ell(w, z = (x, y)) = 0 + 2\lambda w$$

Hence

$$v = \begin{cases} 2\lambda w, & yw^\top x > 1 \\ 2\lambda w, & yw^\top x = 1 \\ -yx + 2\lambda w, & yw^\top x < 1 \end{cases}$$

(4)

Algorithm. For $t = 1, 2, 3, \dots$ iterate:

- Get a random sub-sample $\left\{ \tilde{z}_i^{(t)} = \left(\tilde{x}_i^{(t)}, \tilde{y}_i^{(t)} \right); i = 1, \dots, m \right\}$ of size m with or without replacement from the complete data-set \mathcal{S}_n .
- For $j = 1, \dots, d$ (index j indicates the dimension of w) compute

$$w_j^{(t+1)} = w_j^{(t)} - \eta_t \frac{1}{\sqrt{[G_t]_{j,j} + \epsilon}} \bar{v}_{t,j}$$

$[G_t]_{j,j} = [G_{t-1}]_{j,j} + (\bar{v}_{t,j})^2$ where $\bar{v}_t = \frac{1}{m} \sum_{i=1}^m \tilde{v}_{t,i}$ and

$$\tilde{v}_{t,i} = \begin{cases} 2\lambda w^{(t)}, & \tilde{y}_i^{(t)} (w^{(t)})^\top \tilde{x}_i^{(t)} > 1 \\ 2\lambda w^{(t)}, & \tilde{y}_i^{(t)} (w^{(t)})^\top \tilde{x}_i^{(t)} = 1 \\ -\frac{1}{m} \tilde{y}_i^{(t)} \tilde{x}_i^{(t)} + 2\lambda w^{(t)}, & \tilde{y}_i^{(t)} (w^{(t)})^\top \tilde{x}_i^{(t)} < 1 \end{cases}$$

where index i indicates the sub-sample, and $\epsilon > 0$ small.

- Terminate if a termination criterion is satisfied

(5)

- The R code can be found in the link https://raw.githubusercontent.com/georgios-stats/Machine_Learning_and_Neural_Networks_III_Epiphany_2024/main/Exercises/supplementary/q6_adagrad.R
- The figures are presented below




FIGURE 1. trace plots

- (c) I found $w = (-2.674615, 3.205785)$
 (d) It belongs to -1

Part 3. Support Vector Machines

The following is given as a homework (Formative assessment 3)

Exercise 15. (★★) Consider a training data set $\mathcal{D} = \{z_i = (x_i, y_i)\}_{i=1}^m$. Consider the Soft-SVM Algorithm that requires the solution of the following quadratic minimization problem (in a slightly modified but equivalent form to what we have discussed)

Primal problem:

$$\begin{aligned}
 (7) \quad & (w^*, b^*, \xi^*) = \arg \min_{(w, b, \xi)} \left(\frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^m \xi_i \right) \\
 (8) \quad & \text{subject to: } y_i (\langle w, x_i \rangle + b) \geq 1 - \xi_i, \quad \forall i = 1, \dots, m \\
 (9) \quad & \xi_i \geq 0, \quad \forall i = 1, \dots, m
 \end{aligned}$$

for some user-specified fixed parameter $C > 0$.

- (1) Specify the Lagrangian function L associated to the above primal quadratic minimization problem, where $\{\alpha_i\}$ are the Lagrange coefficients wrt (8), and $\{\beta_i\}$ are the Lagrange coefficients wrt (9). Write down any possible restrictions on the Lagrange coefficients.
- (2) Compute the dual Lagrangian function denoted as \tilde{L} as a function of the Lagrange coefficients and the data points \mathcal{D} .

- (3) Apply the Karush–Kuhn–Tucker (KKT) conditions to the above problem, and write them down.
- (4) Derive and write down the dual Lagrangian quadratic maximization problem, along with the inequality and equality constraints, where you seek to find $\{\alpha_i\}$.
- (5) Justify why the i -th point x_i lies on the margin boundary when $\alpha_i \in (0, C)$ (beware it is $\alpha_i \neq C$), and why the i -th point x_i lies inside the margin when $\alpha_i = C$.
- (6) Given optimal values $\{\alpha_i^*\}$ for Lagrangian coefficients $\{\alpha_i\}$ as they are derived by solving the dual Lagrangian maximization problem in part 4, derive the optimal values w^* and b^* for the parameters w and b as function of the support vectors. Regarding parameter b it should be in the derived in the form

$$b^* = \frac{1}{|\mathcal{M}|} \sum_{i \in \mathcal{M}} \left(y_i - \sum_{j \in \mathcal{S}} \alpha_j^* y_j \langle x_j, x_i \rangle \right)$$

where you determine the sets \mathcal{M} and \mathcal{S} .

- (7) Report the halfspace predictive rule $h_{w,b}(x)$ of the above problem as a function of α^* and b^* .

Solution.

Exercise 16. (★★) *[This is the Relevance Vector Machine. The Exercise is taken from “Exercise Sheet: Bayesian Statistics” of the module “Bayesian Statistics III/IV (MATH3361/4071)” taught in “Michaelmas term 2021”. The supplementary material in the box was mainly provided for the students who had not been introduced to the SVM ideas or the Kernel trick -so it can be skipped. Also, the supplementary material in the box is presented with a statistical (geostatistical modeling) motivation. The exercise requires basic knowledge of Bayesian statistical inference and in particular the use of Bayes theorem for the computation of the posterior as well as basis probability density calculus. However, the exercise is a useful example of extending the SVM ideas to the Bayesian learning setting.]*

Regarding the statistical model: Long story, short (supplementary material)

Consider that we are interested in recovering the mapping

$$x \xrightarrow{\eta} \eta(x)$$

in the sense that $y \in \mathbb{R}$ is the response (output quantity) that depends on $x = (x_1, \dots, x_d) \in \mathcal{X} \subseteq \mathbb{R}^d$ which is the independent variable (input quantity) in a procedure; E.g.:

- y : precipitation in log scale
- $x = (\text{longitude}, \text{latitude})$: geographical coordinates.

Consider a set of observed data $\{(y_i, x_i)\}_{i=1}^n$, which may be contaminated by additive noise of unknown variance; i.e.

$$y_i = \eta(x_i) + \epsilon_i,$$

where $\epsilon_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma^2)$ and $\sigma^2 > 0$ is unknown. We wish to recover $\eta(x)$ by using the Tikhonov regularization on the functional space \mathcal{H} such that

$$(10) \quad \eta = \arg \min_{\forall \tilde{\eta} \in \mathcal{H}} \left\{ \sum_{i=1}^n L(y_i - \tilde{\eta}(x_i)) + \lambda \|\tilde{\eta}\|_{\mathcal{H}}^2 \right\}$$

By assuming that \mathcal{H} is a Reproducing Kernel Hilbert Space (RKHS), the solution to the above Ridge regularizes loss minimization problem is such that

$$\eta(x) = \beta_0 + \sum_{j=1}^n k(x, x_j) \beta_j = k(x)^\top \beta$$

where $k(x) = (1, k(x, x_1), \dots, k(x, x_n))^\top$, $k(x, x_j)$ is the reproducing kernel (such as $k_\phi(x, x_j) = \exp(-\phi \|x - x_j\|^2)$ for some known parameter $\phi > 0$), and $\beta \in \mathbb{R}^{n+1}$ is an unknown vector.

Consider the following Bayesian model²

$$\begin{cases} y|\beta, \sigma^2 & \sim \mathcal{N}(K\beta, I\sigma^2) \\ \beta|\lambda & \sim \mathcal{N}(0, D^{-1}), \quad D = (\lambda_0, \lambda_1, \dots, \lambda_n) \\ \lambda_i & \stackrel{\text{iid}}{\sim} d\Pi(\lambda_i) \propto \lambda_i^{a-1} \exp(-b\lambda_i) d\lambda_i, \quad \forall i = 1, \dots, n \\ \sigma^2 & \sim d\Pi(\sigma^2) \propto (\sigma^2)^{c-1} \exp(-\frac{1}{\sigma^2}d) d\sigma^2 \\ \beta, \sigma^2 & \text{a priori independent} \end{cases}$$

where K is a known matrix with size $n \times (n+1)$ such that

$$K = \begin{bmatrix} 1 & k(x_1, x_1) & \cdots & k(x_1, x_n) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & k(x_n, x_1) & \cdots & k(x_n, x_n) \end{bmatrix}.$$

The quantities $a > 0$, $b > 0$, $c > 0$, $d > 0$, and $\phi > 0$ are considered as fixed.

²Dixit, A., & Roy, V. (2021). Posterior impropriety of some sparse Bayesian learning models. *Statistics & Probability Letters*, 171, 109039.

- (1) When $b = 0$, show that a necessary condition for a valid posterior inference is $a \in (-1/2, 0)$ for any choice of prior for τ (i.e. any choice of (c, d)).
- (2) Let $P = K (K^\top K)^{-1} K^\top$. Show that (2a) and (2b) are sufficient conditions for the Bayesian model to lead to a valid posterior inference
 - (a) if $a > 0$ and $b > 0$, or
 - (b) if $y^\top (I - P) y + 2d > 0$ and $c > -\frac{n}{2}$
- (3) Does the the improper Uniform prior on the joint $\log(\lambda_i)$ and $\log(\sigma^2)$, i.e. $\pi(\log(\lambda_i), \log(\sigma^2)) \propto 1$, lead to a valid inference?
- (4) Does the Jeffreys' prior $\pi(\lambda_i) \propto 1/\lambda_i$ lead to a valid inference?

Hint-1::

$$(y - K\beta)^\top (y - K\beta) + (\beta - \mu)^\top V^{-1}(\beta - \mu) = (\beta - \mu^*)^\top (V^*)^{-1}(\beta - \mu^*) + S^*;$$

$$S^* = \mu^\top V^{-1}\mu - (\mu^*)^\top (V^*)^{-1}(\mu^*) + y^\top y; \quad V^* = (V^{-1} + K^\top K)^{-1}; \quad \mu^* = V^* (V^{-1}\mu + K^\top y)$$

Hint-2:: Sherman-Morrison-Woodbury formula:

$$(A + UCV)^{-1} = A^{-1} - A^{-1}U(C^{-1} + VA^{-1}U)^{-1}VA^{-1}$$

Hint-3::

$$-\frac{y^\top y}{2\sigma^2} \leq -\frac{y^\top (I\sigma^2 + KD^{-1}K^\top)^{-1}y}{2} \leq -\frac{1}{2\sigma^2}y^\top (I - P)y$$

where $P = K (K^\top K)^{-1} K^\top$.

Hint-4:: It is given that $\int_{(0,\infty)} \frac{t^{-(a+1)}}{(\xi+t)^{1/2}} dt < \infty$ if and only if $a \in (-1/2, 0)$.

Solution.

The posterior pdf is given by

$$\pi(\beta, \sigma^2, \lambda|y) = \frac{f(y|\beta, \sigma^2) \pi(\beta, \sigma^2, \lambda)}{f(y)}$$

and is proper iff $f(y) < \infty$ where

$$f(y) = \int \left(\underbrace{\int \left(\underbrace{\int f(y|\beta, \sigma^2) \pi(\beta|\lambda) d\beta}_{=f(y|\lambda, \sigma^2)} \right) \pi(\lambda) d\lambda}_{=f(y|\sigma^2)} \right) \pi(\sigma^2) d\sigma^2$$

It is

$$\begin{aligned}
f(y|\lambda, \sigma^2) &= \int f(y|\beta, \sigma^2) \pi(\beta, \sigma^2) d\beta \\
&= (2\pi)^{-\frac{n+n+1}{2}} \left(\frac{1}{\sigma^2}\right)^{\frac{n}{2}} \det(D)^{\frac{1}{2}} \int \exp\left(-\frac{1}{2\sigma^2} \left((y - K\beta)^\top (y - K\beta) + \beta^\top (D\sigma^2) \beta\right)\right) d\beta \\
&= (2\pi)^{-\frac{n+n+1}{2}} \left(\frac{1}{\sigma^2}\right)^{\frac{1}{2}} \det(D)^{\frac{1}{2}} \left[\int \exp\left(-\frac{1}{2\sigma^2} (\beta - \mu^*)^\top V^* (\beta - \mu^*)\right) d\beta \right] \left[\exp\left(-\frac{1}{2\sigma^2} S^*\right) \right]
\end{aligned}$$

Because

$$\begin{aligned}
\int \exp\left(-\frac{1}{2\sigma^2} (\beta - \mu^*)^\top V^* (\beta - \mu^*)\right) d\beta &= (2\pi)^{\frac{n+1}{2}} \det(V^*/\sigma^2)^{-\frac{1}{2}} \\
&= (2\pi)^{\frac{n+1}{2}} \det(K^\top K + D\sigma^2)^{-\frac{1}{2}}
\end{aligned}$$

$$\begin{aligned}
\exp\left(-\frac{1}{2\sigma^2} S^*\right) &= \exp\left(-\frac{1}{2\sigma^2} \mu^\top (D\sigma^2) \mu - (\mu^*)^\top (V^*)^{-1} (\mu^*) + y^\top y\right) \\
&= \exp\left(-\frac{1}{2\sigma^2} \left(y^\top y - y^\top K (K^\top K + D\sigma^2)^{-1} K^\top y\right)\right) \\
&= \exp\left(-\frac{1}{2\sigma^2} \left(y^\top \left(I - K (K^\top K + D\sigma^2)^{-1} K^\top\right) y\right)\right) \\
&= \exp\left(-\frac{1}{2\sigma^2} \left(y^\top (K^\top D^{-1} K + I\sigma^2)^{-1} y\right)\right)
\end{aligned}$$

So

$$\begin{aligned}
f(y|\lambda, \sigma^2) &= (2\pi)^{-\frac{n}{2}} \left(\frac{1}{\sigma^2}\right)^{\frac{1}{2}} \det(D)^{\frac{1}{2}} \det(K^\top K + \sigma^2 D)^{-\frac{1}{2}} \\
&\quad \times \exp\left(-\frac{1}{2\sigma^2} \left(y^\top (I\sigma^2 + K^\top D^{-1} K)^{-1} y\right)\right)
\end{aligned}$$

(1) I have

$$\begin{aligned}
f(y|\sigma^2) &= \int f(y|\lambda, \sigma^2) \pi(\lambda) d\lambda \\
&= (2\pi)^{-\frac{n}{2}} (\sigma^2)^{\frac{1}{2}} \int \left[\det(D)^{\frac{1}{2}} \right] \left[\det(K^\top K + D\sigma^2)^{-\frac{1}{2}} \right] \\
&\quad \times \exp\left(-\frac{1}{2} \left(y^\top (I\sigma^2 + K^\top D^{-1} K)^{-1} y\right)\right) \left[\prod_{i=0}^n \lambda_i^{a-1} \right] d\lambda_0 \dots d\lambda_n
\end{aligned}$$

because $b = 0$.

- It is $\exp\left(-\frac{y^\top y}{2\sigma^2}\right) \leq \exp\left(-\frac{y^\top (I\sigma^2 + K^\top D^{-1} K)^{-1} y}{2}\right)$
- It is $\det(D)^{\frac{1}{2}} = \prod_{i=0}^n \lambda_i^{\frac{1}{2}}$.
- If $\{e_j\}_{j=0}^{n-1}$ are eigenvalues of $K^\top K$ and $e_{\max} = \max(\{e_j\})$, then $K^\top K + D\sigma^2 \leq Ie_{\max} + D\sigma^2$, consequently $\det(K^\top K + D\sigma^2)^{-\frac{1}{2}} \geq \prod_{j=0}^n (\lambda_j \sigma^2 + e_{\max})^{-\frac{1}{2}}$.

Then

$$\begin{aligned}
f(y|\sigma^2) &\geq (2\pi)^{-\frac{n}{2}} (\sigma^2)^{\frac{1}{2}} \int \prod_{j=0}^n \lambda_j^{\frac{1}{2}} \prod_{j=0}^n (\lambda_j \sigma^2 + e_{\max})^{-\frac{1}{2}} \exp\left(-\frac{1}{2\sigma^2} y^\top y\right) \prod_{j=0}^n \lambda_j^{a-1} d\lambda_0 \dots d\lambda_n \\
&= (2\pi)^{-\frac{n}{2}} (\sigma^2)^{\frac{1}{2}} \exp\left(-\frac{1}{2\sigma^2} y^\top y\right) \int \dots \int \prod_{j=0}^n \left[\lambda_j^{\frac{1}{2}}\right] \left[\prod_{j=0}^n (\lambda_j \sigma^2 + e_{\max})^{-\frac{1}{2}}\right] \left[\prod_{j=0}^n \lambda_j^{a-1}\right] d\lambda_0 \dots d\lambda_n \\
&= (2\pi)^{-\frac{n}{2}} (\sigma^2)^{\frac{1}{2}} \exp\left(-\frac{1}{2\sigma^2} y^\top y\right) \prod_{j=0}^n \int \frac{\lambda_j^{a-\frac{1}{2}}}{(\lambda_j \sigma^2 + e_{\max})^{\frac{1}{2}}} d\lambda_j
\end{aligned}$$

Let $t_i = 1/\lambda_i$, then

$$f(y|\sigma^2) \geq (2\pi)^{-\frac{n}{2}} (\sigma^2)^{\frac{1}{2}} \exp\left(-\frac{1}{2\sigma^2} y^\top y\right) \prod_{j=0}^n \int \frac{t_{jj}^{-a-1}}{\left(t_j + \frac{\sigma^2}{e_{\max}}\right)^{\frac{1}{2}}} d\lambda_j$$

which is finite if and only if $a \in (-1/2, 0)$.

(2)

- (a) If $a > 0$, $b > 0$ then $\lambda_i \stackrel{\text{iid}}{\sim} \text{Ga}(a, b)$ for all $i = 1, \dots, n$, and if $c > 0$, $d > 0$ then $\tau \stackrel{\text{iid}}{\sim} \text{Ga}(c, d)$ which are proper. So $\Pi(\beta, \sigma^2, \lambda, \tau)$ is a proper prior, and hence it leads to proper posterior.
- (b) I have

$$\begin{aligned}
f(y|\sigma^2) &= (2\pi)^{-\frac{n}{2}} (\sigma^2)^{\frac{1}{2}} \int \left[\det(D)^{\frac{1}{2}}\right] \left[\det(K^\top K + D\sigma^2)^{-\frac{1}{2}}\right] \\
&\quad \times \exp\left(-\frac{1}{2} \left(y^\top (I\sigma^2 + K^\top D^{-1}K)^{-1} y\right)\right) \pi(\lambda) d\lambda
\end{aligned}$$

It is $\det(D)^{\frac{1}{2}} = \prod_{i=0}^n \lambda_i^{\frac{1}{2}}$. Also, it is $K^\top K + D\sigma^2 \geq D\sigma^2$ then $\det(K^\top K + D\sigma^2)^{-\frac{1}{2}} \leq \prod_{j=0}^n (\lambda_j \sigma^2)^{-\frac{1}{2}}$. Hence

$$f(y|\sigma^2) \leq (2\pi)^{-\frac{n}{2}} (\sigma^2)^{\frac{n}{2}} \exp\left(-\frac{1}{2\sigma^2} y^\top (I - P) y\right) \int \pi(\lambda) d\lambda$$

which implies that $f(y|\sigma^2) < \infty$ if $\pi(\lambda)$ is proper. Yet,

$$\begin{aligned}
f(y) &= \int f(y|\sigma^2) \pi(\sigma^2) d\sigma^2 \\
&\leq (2\pi)^{-\frac{n}{2}} \int (\sigma^2)^{-\frac{n}{2}+c+1} \exp\left(-\frac{1}{\sigma^2} \left(\frac{y^\top (I - P) y}{2} + d\right)\right) d\sigma^2
\end{aligned}$$

which is finite if $y^\top (I - P) y + 2d > 0$ and $c > -\frac{n}{2}$.

- (c) No. This implies $\pi(\lambda, \sigma^2) \propto \sigma^2 \prod_{j=0}^n \lambda_j^{-1}$. It is improper prior as $\int \pi(\lambda, \sigma^2) d(\lambda, \sigma^2) = \infty$, and $(a, b, c, d) = (0, 0, 0, 0)$ which violates the necessary conditions.
- (d) No, it violates the necessary conditions.

Exercise 17. (★) Students are encouraged to practice on the Exercises 6.1-6.19 from the textbook

- Bishop, C. M. (2006). Pattern recognition and machine learning (Vol. 4, No. 4, p. 738). New York: Springer.

available from

- <https://www.microsoft.com/en-us/research/uploads/prod/2006/01/Bishop-Pattern-Recognition-and-Machine-Learning-2006.pdf>

The solutions are available from

- https://blackboard.durham.ac.uk/ultra/courses/_44662_1/outline/create/document?id=_1396738_1
-

Part 4. The kernel trick

Exercise 18. (★★) Show that K with

$$K(x, y) = \frac{\sin\left(2\pi\left(N + \frac{1}{2}\right)(x - y)\right)}{\sin(\pi(x - y))}$$

is a valid kernel.

Hint-1: You may use that $\sum_{n=0}^r z^n = \frac{1-z^{r+1}}{1-z}$

Hint-2: You may use that $e^{ix} = \cos(x) + i \sin(x)$

Solution. It is

$$\begin{aligned} K(x, y) &= \frac{\sin\left(2\pi\left(N + \frac{1}{2}\right)(x - y)\right)}{\sin(\pi(x - y))} = \frac{-2i \sin\left(2\pi\left(N + \frac{1}{2}\right)(x - y)\right)}{-2i \sin(\pi(x - y))} \\ &= \frac{e^{-2\pi\left(N + \frac{1}{2}\right)i(x-y)} - e^{2\pi\left(N + \frac{1}{2}\right)i(x-y)}}{e^{-\pi i(x-y)} - e^{\pi i(x-y)}} \\ &= \frac{e^{\pi i(x-y)} e^{-2\pi\left(N + \frac{1}{2}\right)i(x-y)} - e^{2\pi\left(N + \frac{1}{2}\right)i(x-y)}}{e^{\pi i(x-y)} e^{-\pi i(x-y)} - e^{\pi i(x-y)}} \\ &= e^{-2\pi i N(x-y)} \frac{1 - (e^{2\pi i(x-y)})^{2N+1}}{1 - e^{2\pi i(x-y)}} \\ &= e^{-2\pi i N(x-y)} \sum_{n=0}^{2N} \left(e^{2\pi i(x-y)}\right)^n = \sum_{n=-N}^N e^{2\pi i n(x-y)} = \sum_{n=-N}^N e^{2\pi i n x} e^{-2\pi i n y} \\ &= \sum_{n=-N}^N e^{2\pi i n x} \overline{e^{2\pi i n y}} = \langle \psi(x), \psi(y) \rangle \end{aligned}$$

with $\psi(x) = (e^{-2\pi i N x}, e^{-2\pi i (N-1)x}, \dots, 1, \dots, e^{2\pi i (N-1)x}, e^{2\pi i N x})^\top$. Based on the theorem in the Handout, the Kernel can be expressed as an inner product of a vector of bases, hence it is a valid kernel.

Note that given Hint 2 it is

$$\begin{aligned}
e^{-2\pi i n y} &= \cos(-2\pi n y) + i \sin(-2\pi n y) \\
&= \cos(2\pi n y) - i \sin(2\pi n y) \\
&= \overline{\cos(2\pi n y) + i \sin(2\pi n y)} \\
&= \overline{e^{2\pi i n y}}
\end{aligned}$$

Exercise 19. (★★) (Kernel ridge regression) Consider the standard linear regression problem with learning rule $h_\theta(x) = \theta^\top x$. Consider a training data set $\mathcal{S} = \{z_i = (x_i, y_i)\}_{i=1}^m$. The ridge regression cost function is then

$$\hat{R}_{\mathcal{S}}(\theta) = \frac{1}{2} \sum_{i=1}^m \left(\theta^\top x_i - y_i \right)^2 + \frac{\lambda}{2} \|\theta\|_2^2$$

- (1) Use a vector notation to find a closed-form expression for the value of θ which minimizes the ridge regression cost function.
- (2) Suppose that we want to use kernels to implicitly represent our feature vectors in a high-dimensional (possibly infinite dimensional) space. Using a feature mapping ψ , the ridge regression cost function becomes

$$\hat{R}_{\mathcal{S}}^{\psi}(\theta) = \frac{1}{2} \sum_{i=1}^m \left(\theta^\top \psi(x_i) - y_i \right)^2 + \frac{\lambda}{2} \|\theta\|_2^2$$

Making a prediction on a new input x_{new} would now be done by computing $\theta^\top \psi(x_{\text{new}})$. Show how we can use the “kernel trick” to obtain a closed form for the prediction on the new input x_{new} without ever explicitly computing $\psi(x_{\text{new}})$. You may assume that the parameter vector θ can be expressed as a linear combination of the input feature vectors; i.e., $\theta = \sum_{i=1}^m \alpha_i \psi(x_i)$ for some set of parameters $\{\alpha_i\}$.

Hint:: You may need to use the identity

$$(\lambda I + BA)^{-1} B = B(\lambda I + AB)^{-1}.$$

Solution.

- (1) It is

$$\hat{R}_{\mathcal{S}}(\theta) = \frac{1}{2} (X\theta - y)^\top (X\theta - y) + \frac{\lambda}{2} \theta^\top \theta$$

The gradient is

$$\nabla_{\theta} \hat{R}_{\mathcal{S}}(\theta) = X^\top X\theta - X^\top y + \lambda\theta$$

and by setting it equal to zero at $\hat{\theta}$, I get

$$\begin{aligned}
0 &= \nabla_{\theta} \hat{R}_{\mathcal{S}}(\theta) \Big|_{\theta=\hat{\theta}} = X^\top X\hat{\theta} - X^\top y + \lambda\hat{\theta} \\
&\implies \hat{\theta} = \left(X^\top X - \lambda I \right)^{-1} X^\top y
\end{aligned}$$

- (2) Now the design matrix is Ψ associated with the feature vectors $\psi(x_i)$ i.e. $[\Psi]_{i,j} = \psi_j(x_i)$. Then from the previous part, it is

$$\begin{aligned}\hat{\theta} &= \left(\Psi^\top \Psi - \lambda I \right)^{-1} \Psi^\top y \\ &= \Psi^\top \left(\Psi^\top \Psi - \lambda I \right)^{-1} y \\ &= \Psi^\top (K - \lambda I)^{-1} y\end{aligned}$$

where K is the Gram for the training set i.e. $[K]_{i,j} = k(x_i, x_j)$ there $k(\cdot, \cdot)$ is a kernel such as $k(a, b) = \psi(a)^\top \psi(b)$. To predict a new value $y_{\text{new}} = h_\theta(x_{\text{new}})$, we can compute

$$\begin{aligned}y_{\text{new}} &= h_{\hat{\theta}}(x_{\text{new}}) = \hat{\theta}^\top \psi(x_{\text{new}}) \\ &= \left(\Psi^\top (K - \lambda I)^{-1} y \right)^\top \psi(x_{\text{new}}) \\ &= y^\top (K - \lambda I)^{-1} \Psi \psi(x_{\text{new}}) \\ &= \sum_{i=1}^m \alpha_i K \psi(x_i, x_{\text{new}})\end{aligned}$$

where $\alpha_i = (K - \lambda I)^{-1} y$

Exercise 20. (★★) (On the SVM with Gaussian kernel) Consider the task of training a support vector machine using the Gaussian kernel

$$k(a, b) = \exp \left(-\frac{1}{\phi^2} \|a - b\|_2^2 \right)$$

We will show that as long as there are no two identical points in the training set, we can always find a value for the scale parameter $\phi > 0$ such that the SVM achieves zero training error. Consider a training data set $\mathcal{S} = \{z_i = (x_i, y_i)\}_{i=1}^m$.

- (1) Recall from class that the decision function learned by the support vector machine can be written as

$$h_w(x) = \sum_{i=1}^m \alpha_i y_i k(x_i, x) + b$$

Assume that the training data $\mathcal{S} = \{z_i = (x_i, y_i)\}_{i=1}^m$ consists of points which are separated by at least a distance of ϵ that is $\|x_i - x_j\|_2 \geq \epsilon$ for any $i \neq j$. Find values for the set of parameters $\{\alpha_i\}$ and b and Gaussian kernel scale parameter ϕ such that x_i is correctly classified, for all $i = 1, \dots, m$.

Hint:: Let $\alpha_i = 1$ for all i and $b = 0$. Now notice that for $y \in \{-1, +1\}$ the prediction on x_i will be correct if $|f(x_i) - y_i| < 1$, so find a value of ϕ that satisfies this inequality for all i .

- (2) (Requires additional reading) Suppose we run a SVM with slack variables using the parameter $\phi > 0$ you found in the previous part. Consider that you are using the “Sequential

minimal optimization” to solve the quadratic programming (QP) problem that arises during the training of support-vector machines (SVM) –for details see:

https://en.wikipedia.org/wiki/Sequential_minimal_optimization

Will the resulting classifier necessarily obtain zero training error? Why or why not? A short explanation (without proof) will suffice.

Solution.

- (1) First we set $\alpha_i = 1$ for all $i = 1, \dots, m$ and $b = 0$. Then, for a training example $\{x_i, y_i\}$, we get

$$\begin{aligned}
 |h_w(x_i) - y_i| &= \left| \sum_{j=1}^m y_j k(x_j, x_i) - y_i \right| \\
 &= \left| \sum_{j=1}^m y_j \exp\left(-\frac{1}{\phi^2} \|x_j - x_i\|_2^2\right) - y_i \right| \\
 &= \left| y_i + \sum_{j \neq i} y_j \exp\left(-\frac{1}{\phi^2} \|x_j - x_i\|_2^2\right) - y_i \right| \\
 &\leq \sum_{j \neq i} \left| y_j \exp\left(-\frac{1}{\phi^2} \|x_j - x_i\|_2^2\right) \right| \\
 &= \sum_{j \neq i} |y_j| \exp\left(-\frac{1}{\phi^2} \|x_j - x_i\|_2^2\right) \\
 &\leq \sum_{j \neq i} |y_j| \exp\left(-\frac{1}{\phi^2} \epsilon^2\right) \\
 &= (m-1) \exp\left(-\frac{1}{\phi^2} \epsilon^2\right)
 \end{aligned}$$

Thus we need to choose a γ such that

$$(m-1) \exp\left(-\frac{1}{\phi^2} \epsilon^2\right) < 1 \iff \phi < \frac{\epsilon}{\log(m-1)}$$

E.g. $\phi = \frac{\epsilon}{\log(m)}$

- (2) The classifier will obtain zero training error. The SVM without slack variables will always return zero training error if it is able to find a solution, so all that remains to be shown is that there exists at least one feasible point. Consider the constraint $y_i(w^\top x_i + b)$ for some i , and let $b = 0$. Then

$$y_i (w^\top x_i + b) = y_i f(x_i) > 0$$

since $f(x_i)$ and y_i have the same sign, and shown above. Therefore, as we choose all the α_i 's large enough, $y_i (w^\top x_i + b) > 1$, so the optimization problem is feasible.

Exercise 21. (★★) (Kernel Principal Component Analysis) Consider the standard PCA. Consider a given dataset in the form of a matrix X such as $[X]_{i,j} = x_{i,j}$ for $i = 1, \dots, m$ and $j = 1, \dots, p$, and $x_i = (x_{i,1}, \dots, x_{i,p})^\top$. Assume the data are centered around zero as $\sum_{i=1}^m x_i = 0$. The Sample covariance matrix is

$$C = \frac{1}{m} \sum_{i=1}^m x_i x_i^\top$$

Recall that the i th Principal Component is $y_i = v_j^\top x_i$ where V is a $p \times p$ matrix of weights whose columns are the eigenvectors of $X^\top X$; i.e. $X^\top X = V^\top \Lambda V$, and $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_p)$.

Suppose that we want to use kernels to implicitly represent our feature vectors in a high-dimensional (possibly infinite dimensional) space. Using a feature mapping ψ , the Sample covariance matrix is

$$C^\psi = \frac{1}{m} \sum_{i=1}^m \psi(x_i) (\psi(x_i))^\top$$

Show how we can use the “kernel trick” to obtain a closed form for the Principal Components without ever explicitly computing $\psi(\cdot)$. You can focus on the first principal component only.

Solution.

The Sample covariance matrix is $C^\psi = \frac{1}{m} \sum_{i=1}^m \psi(x_i) (\psi(x_i))^\top$, and keep on assuming that the transformation keeps our data centered around zero. To find the first PC we need to solve

$$C^\psi v_1 = \lambda_1 v_1 \iff \frac{1}{m} \sum_{i=1}^m \psi(x_i) (\psi(x_i))^\top v_1 = \lambda_1 v_1$$

Since $\lambda_1 \neq 0$, v_1 must be the span of $\psi(x_i)$ i.e. v_1 can be written as

$$v_1 = \sum_j \alpha_j \psi(x_j)$$

So by substituting, I get

$$\frac{1}{m} \sum_{i=1}^m \psi(x_i) (\psi(x_i))^\top \sum_j \alpha_j \psi(x_j) = \lambda_1 \sum_j \alpha_j \psi(x_j)$$

To get only inner products, I multiply both sides with $\psi(x_n)$ for any $n = 1, \dots, m$, and I get

$$\begin{aligned} \frac{1}{m} \sum_{i=1}^m \psi(x_i) (\psi(x_i))^\top \sum_j \alpha_j (\psi(x_n))^\top \psi(x_j) &= \lambda_1 \sum_j \alpha_j (\psi(x_n))^\top \psi(x_j) \iff \\ \frac{1}{m} \sum_j \alpha_j \sum_{i=1}^m [(\psi(x_n))^\top \psi(x_i)] [(\psi(x_i))^\top \psi(x_j)] &= \lambda_1 \sum_j \alpha_j [(\psi(x_n))^\top \psi(x_j)] \end{aligned}$$

Now we can set/specify a kernel functions as

$$k(a, b) = (\psi(a))^\top \psi(b)$$

and substitute the inner product with the above kernel as

$$\frac{1}{m} \sum_j \alpha_j \sum_{i=1}^m k(x_n, x_i) k(x_i, x_j) = \lambda_1 \sum_j \alpha_j k(x_n, x_j)$$

which allow us to solve the problem without working directly in a high-dimensional space. Let K be a Gram matrix such as $[K]_{i,j} = k(x_i, x_j)$. Hence I get

$$\frac{1}{m} K^\top K \alpha = \lambda_1 K \alpha$$

or for $\tilde{\alpha} = K \alpha$ and $\tilde{\lambda}_1 = m \lambda_1$

$$K \tilde{\alpha} = \tilde{\lambda}_1 \tilde{\alpha}$$

Hence we have to simply solve the eigenvalue problem on the Gram matrix. For the rest principal components, we follow the same procedure.

To extract the first principal component we simply take

$$y_i = \sum_j \alpha_j (\psi(x_j))^\top \psi(x_i) = \sum_j \alpha_j k(x_i, x_j)$$

for each i .

Part 5. Artificial Neural Networks

Exercise 22. (★) Consider the regression problem, with a predictive rule $h_w : \mathbb{R}^d \rightarrow \mathbb{R}$, as a classification probability, that receives values $x \in \mathbb{R}^d$ returns values in \mathbb{R} . Let $h_w(x)$ be modeled as an ANN

$$h(x) = \sigma_2 \left(\sum_{j=1}^c w_{2,1,j} \sigma_1 \left(\sum_{i=1}^d w_{1,j,i} x_i \right) \right)$$

and let the associated activation function be

$$\sigma_2(a) = a \Phi(a)$$

where $\Phi(x) = \int_{-\infty}^x \phi(t) dt$ is considered as known function, and $\phi(t) = \frac{1}{\sqrt{2\pi}} \exp(-\frac{1}{2}t^2)$ and

$$\sigma_1(a) = \exp(-a^2)$$

Consider a loss

$$\ell(w, z = (x, y)) = \frac{1}{2} (y - h_w(x))^2$$

at w and example $z = (x, y)$, where $x \in \mathbb{R}^d$ is the input vector (features), and y is the output vector (targets) with $y \in \mathbb{R}$. Consider that d , c , and q are known integers.

- (1) Perform the forward pass of the back-propagation procedure to compute the activations which may be denoted as $\{a_{t,i}\}$ and outputs which may be denoted as $\{o_{t,i}\}$ at each layer t .
- (2) Perform the backward pass of the back-propagation procedure in order to compute the elements of the gradient $\nabla_w \ell(w, (x, y))$.

Solution.

(1)

Forward pass:

Set: $o_{0,i} = x_i$ for $i = 1, \dots, d$

Compute:

at $t = 1$: **for** $j = 1, \dots, c$

comp: $\alpha_{1,j} = \sum_{i=1}^d w_{1,j,i} x_i$

comp: $o_{1,j} = \exp(-\alpha_{1,j}^2)$

at $t = 2$:

comp: $\alpha_{2,1} = \sum_{j=1}^c w_{2,1,j} o_{1,j}$

comp: $o_{2,1} = \alpha_{2,1} \Phi(\alpha_{2,1})$

get: $h_1 = o_{2,1}$

(2) It is

$$\frac{d}{da} \sigma_1(a) = -2a \exp(-a^2)$$

and

$$\frac{d}{da} \sigma_2(a) = \Phi(a) + a\phi(a)$$

and

$$\frac{d\ell_2}{d\alpha_{2,1}} = \frac{d\ell_2}{do_{2,1}} \frac{do_{2,1}}{d\alpha_{2,1}} = (o_{2,1} - y_1) (\Phi(\alpha_{2,1}) + \alpha_{2,1}\phi(\alpha_{2,1}))$$

Backward pass:

at $t = 2$:

comp: $\tilde{\delta}_{2,1} = \frac{d}{d\alpha_{2,1}} \ell_T = (o_{2,1} - y_1) (\Phi(\alpha_{2,1}) + \alpha_{2,1}\phi(\alpha_{2,1}))$

at $t = 1$: **for** $j = 1, \dots, c$

comp:

$$\begin{aligned} \tilde{\delta}_{1,j} &= \left. \frac{d}{d\xi} \sigma_1(\xi) \right|_{\xi=\alpha_{1,j}} w_{2,1,j} \tilde{\delta}_{2,1} \\ &= -2\alpha_{1,j} \exp(-\alpha_{1,j}^2) w_{2,1,j} \tilde{\delta}_{2,1} \end{aligned}$$

Output:

$$\frac{d}{dw_{1,j,i}} \ell = \tilde{\delta}_{1,j} x_i \text{ and } \frac{d}{dw_{2,1,j}} \ell = \tilde{\delta}_{2,1} o_{1,j}$$

The following is given as a homework (Formative assessment 4)

Exercise 23. (★) Consider the multi-class classification problem, with a predictive rule $h_w : \mathbb{R}^d \rightarrow \mathcal{P}$, as a classification probability i.e, $h_{w,k}(x) = \Pr(x \text{ belongs to class } k)$, that receives values $x \in \mathbb{R}^d$ returns vales in $\mathcal{P} = \left\{ p \in (0, 1)^q : \sum_{j=1}^q p_j = 1 \right\}$. Let $h_w = (h_{w,1}, \dots, h_{w,q})^\top$, let $h_w(x)$ be modeled as an ANN

$$h_k(x) = \sigma_2 \left(\sum_{j=1}^c w_{2,k,j} \sigma_1 \left(\sum_{i=1}^d w_{1,j,i} x_i \right) \right)$$

for $k = 1, \dots, q$, and let the associated activation functions be

$$\sigma_2(a_k) = \frac{\exp(a_k)}{\sum_{k'=1}^q \exp(a_{k'})}, \text{ for } k = 1, \dots, q$$

(called softmax function) and $\sigma_1(a) = \arctan(a)$. Consider a loss

$$\ell(w, z = (x, y)) = - \sum_{k=1}^q y_k \log(h_{w,k}(x))$$

at w and example $z = (x, y)$, where $x \in \mathbb{R}^d$ is the input vector (features), and $y = (y_1, \dots, y_q)$ is the output vector (labels) with $y \in \{0, 1\}^q$ and $\sum_{k=1}^q y_k = 1$. Consider that d , c , and q are known integers.

Hint: You may use

$$\frac{d}{dx} \arctan(x) = \frac{1}{1+x^2}$$

- (1) Perform the forward pass of the back-propagation procedure to compute the activations which may be denoted as $\{a_{t,i}\}$ and outputs which may be denoted as $\{o_{t,i}\}$ at each layer t .
- (2) Show that

$$\frac{\partial}{\partial a_k} \sigma_2(a_j) = \sigma_2(a_j) (1(j=k) - \sigma_2(a_k))$$

$$\text{for } k = 1, \dots, q. \text{ Let } 1(j=k) = \begin{cases} 1 & j = k \\ 0 & j \neq k \end{cases}.$$

- (3) Perform the backward pass of the back-propagation procedure in order to compute the elements of the gradient $\nabla_w \ell(w, (x, y))$.

Solution.

Exercise 24. (★) Students are encouraged to practice on the Exercises 5.1-5.28 from the textbook

- Bishop, C. M. (2006). Pattern recognition and machine learning (Vol. 4, No. 4, p. 738). New York: Springer.

available from

- <https://www.microsoft.com/en-us/research/uploads/prod/2006/01/Bishop-Pattern-Recognition-and-Machine-Learning-2006.pdf>

The solutions are available from

- https://blackboard.durham.ac.uk/ultra/courses/_44662_1/outline/create/document?id=_1396738_1